

# Introduction to Optimization for Statistical Model Fitting

Luigi Acerbi

Department of Computer Science  
University of Helsinki  
Finnish Center for Artificial Intelligence FCAI



BAMB! Summer School – Day 2  
September 2022

## 1 Introduction and recap

- Of models and likelihoods
- The psychometric function

## 2 Model fitting

- A statistical estimation problem
- Model fitting via point estimation

## 3 Intro to optimization

- The problem
- Optimization algorithms
- Optimization cheat sheets

# Learning objectives

By the end of this lecture/tutorial, we will:

- Explain how model fitting is an **optimization problem**
- Describe several major **optimization algorithms**
- Set up and run optimization on **real dataset and models**
- Summarize practical **tips & tricks** for optimization

## 1 Introduction and recap

- Of models and likelihoods
- The psychometric function

## 2 Model fitting

- A statistical estimation problem
- Model fitting via point estimation

## 3 Intro to optimization

- The problem
- Optimization algorithms
- Optimization cheat sheets

# What is a model?



*The best material model of a cat is another, or preferably the same, cat.*

Wiener, *Philosophy of Science* (1945) (with Rosenblueth)

# What is a mathematical model?

- Quantitative stand-in for a theory

# What is a mathematical model?

- Quantitative stand-in for a theory
- A *family of probability distributions* over possible datasets:

$$p(\text{data}|\theta)$$

- ▶ data is a dataset with  $n$  data points (e.g., trials)
- ▶  $\theta$  is a parameter vector

# What is a mathematical model?

- Quantitative stand-in for a theory
- A *family of probability distributions* over possible datasets:

$$p(\text{data}|\theta)$$

- ▶ data is a dataset with  $n$  data points (e.g., trials)
  - ▶  $\theta$  is a parameter vector
- 
- Why?

# What is a mathematical model?

- Quantitative stand-in for a theory
- A *family of probability distributions* over possible datasets:

$$p(\text{data}|\theta)$$

- ▶ data is a dataset with  $n$  data points (e.g., trials)
  - ▶  $\theta$  is a parameter vector
- 
- **Why?** Description, prediction, and explanation

# What is a mathematical model?

- Quantitative stand-in for a theory
- A *family of probability distributions* over possible datasets:

$$p(\text{data}|\theta)$$

- ▶ data is a dataset with  $n$  data points (e.g., trials)
- ▶  $\theta$  is a parameter vector
- **Why?** Description, prediction, and explanation
- Defining  $p(\text{data}|\theta)$  is the core of model building

# What is a mathematical model?

- Quantitative stand-in for a theory
- A *family of probability distributions* over possible datasets:

$$p(\text{data}|\theta)$$

- ▶ data is a dataset with  $n$  data points (e.g., trials)
- ▶  $\theta$  is a parameter vector
- **Why?** Description, prediction, and explanation
- Defining  $p(\text{data}|\theta)$  is the core of model building
  - ▶ Wait, what?

# What is a mathematical model?

- Quantitative stand-in for a theory
- A *family of probability distributions* over possible datasets:

$$p(\text{data}|\theta)$$

- ▶ data is a dataset with  $n$  data points (e.g., trials)
- ▶  $\theta$  is a parameter vector
- **Why?** Description, prediction, and explanation
- Defining  $p(\text{data}|\theta)$  is the core of model building
  - ▶ Wait, what?
- **How?** Think about the data generation process!

# What is a mathematical model?

- Quantitative stand-in for a theory
- A *family of probability distributions* over possible datasets:

$$p(\text{data}|\theta)$$

- ▶ data is a dataset with  $n$  data points (e.g., trials)
- ▶  $\theta$  is a parameter vector
- **Why?** Description, prediction, and explanation
- Defining  $p(\text{data}|\theta)$  is the core of model building
  - ▶ Wait, what?
- **How?** Think about the data generation process!

We need some data

# Data from International Brain Laboratory (IBL)



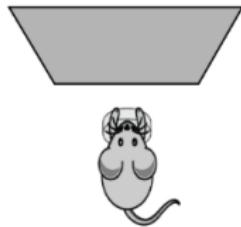
HOME    PUBLICATIONS    RESOURCES    ABOUT    OUR TEAM    JOIN US    IBL MEMBER LOGIN

# International Brain Laboratory

Experimental & theoretical neuroscientists collaborating to understand  
brainwide circuits for complex behavior

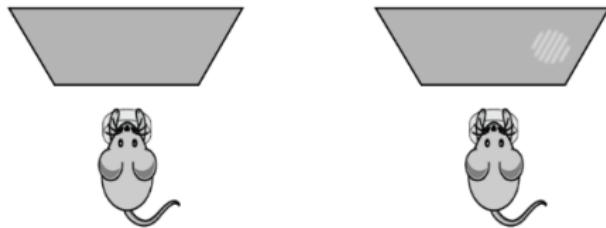
<https://www.internationalbrainlab.com>

# IBL Task



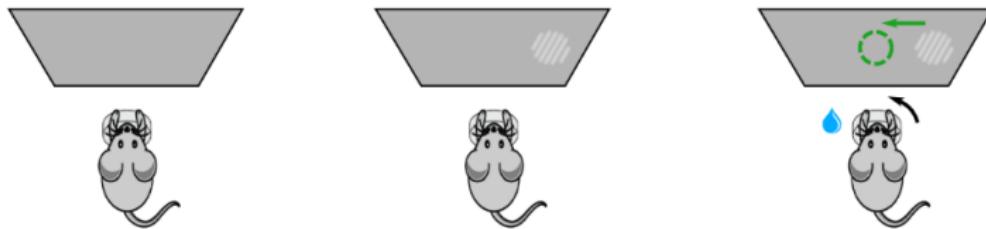
(IBL et al., *eLife*, 2021)

# IBL Task



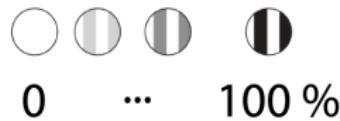
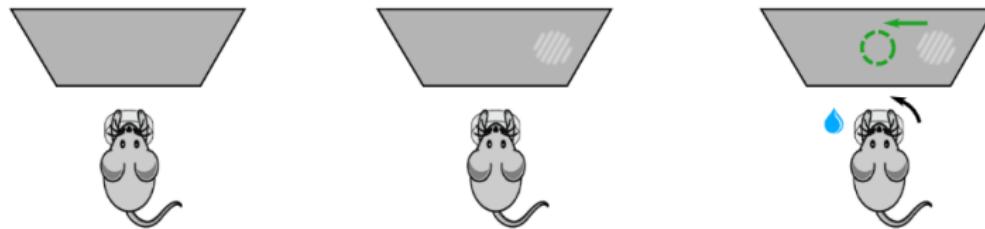
(IBL et al., *eLife*, 2021)

# IBL Task



(IBL et al., *eLife*, 2021)

# IBL Task



(IBL et al., *eLife*, 2021)

# IBL Task

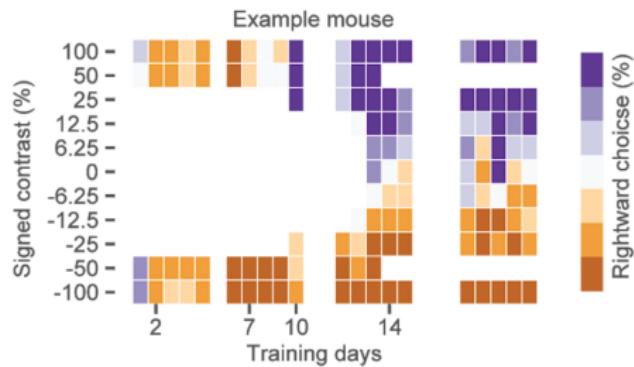


0

...



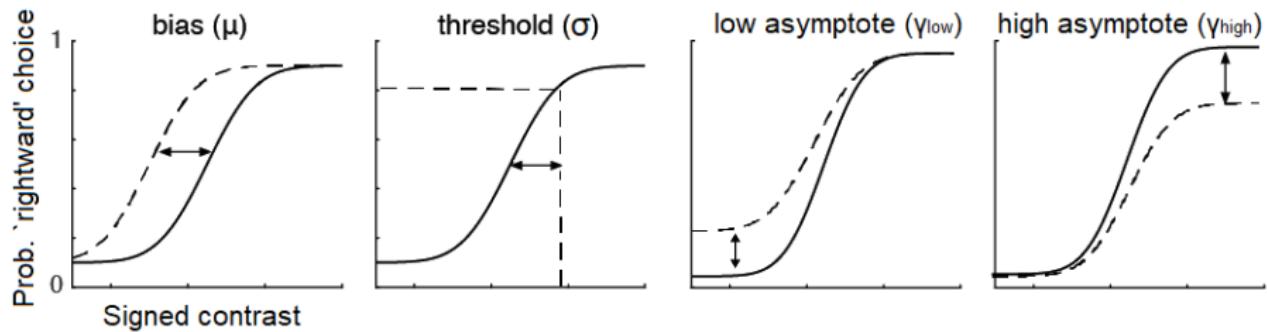
100 %

(IBL et al., *eLife*, 2021)

# Hacking time I

Let's have a look at the data

# The psychometric function



- Data: (signed contrast, choice) for each trial
- Parameters  $\theta$ :  $(\mu, \sigma, \gamma_{\text{low}}, \gamma_{\text{high}})$

$$p(\text{rightward choice} | s, \theta) = \gamma_{\text{low}} + (1 - \gamma_{\text{low}} - \gamma_{\text{high}}) \cdot F(s; \mu, \sigma)$$

# The psychometric function (alt version)

- Default decision process  $F(s; \mu, \sigma)$
- Lapses with probability  $\lambda \in [0, 1]$  (*lapse rate*)
- If lapse, respond 'rightward' with probability  $\gamma \in [0, 1]$  (*lapse bias*)
- Parameters  $\theta$ :  $(\mu, \sigma, \lambda, \gamma)$

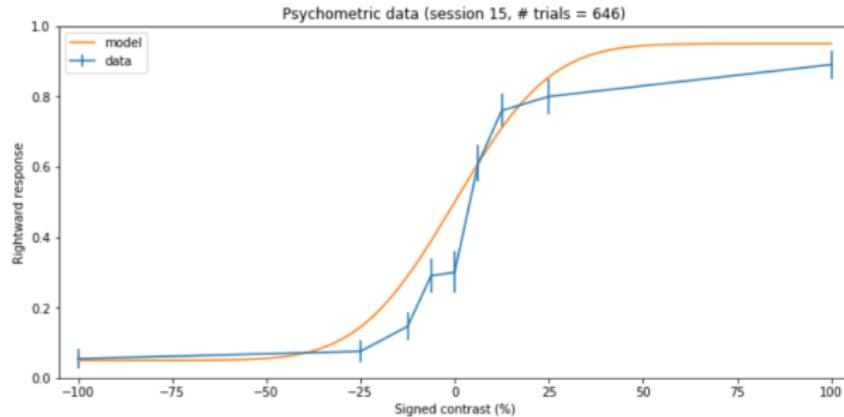
$$p(\text{rightward choice}|s, \theta) = \lambda\gamma + (1 - \lambda) \cdot F(s; \mu, \sigma)$$

## Hacking time II

Let's have a look at the psychometric function

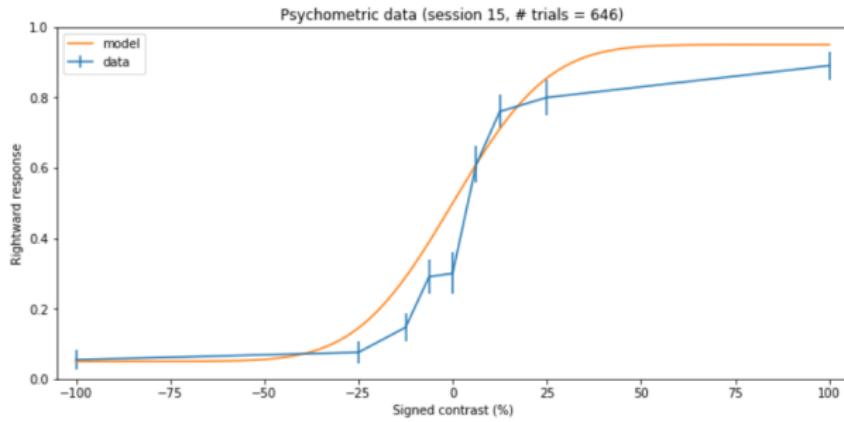
# Metric for model fitting

We need a quantity to measure *goodness of fit*



# Metric for model fitting

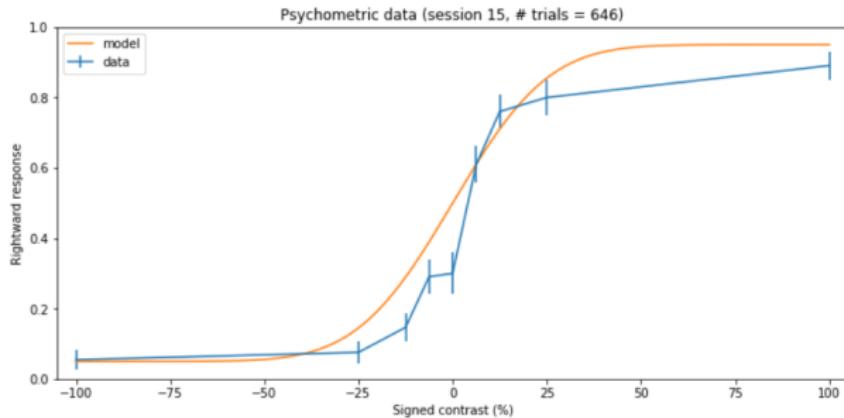
We need a quantity to measure *goodness of fit*



- Mean squared error?

# Metric for model fitting

We need a quantity to measure *goodness of fit*



- Mean squared error?
- The likelihood  $p(\text{data}|\theta)$

## The (log) likelihood

- $p(\text{data}|\theta)$  is a *probability density* as you vary data for a fixed  $\theta$
- $p(\text{data}|\theta)$  is the *likelihood*, a function of  $\theta$  for fixed data

# The (log) likelihood

- For numerical reasons we work with  $\log p(\text{data}|\theta)$

# The (log) likelihood

- For numerical reasons we work with  $\log p(\text{data}|\theta)$
- Using the rules of probability and logarithms:

$$\begin{aligned}\log p(\text{data}|\theta) &= \log p(r^{(1)}, \dots, r^{(n)} | s^{(1)}, \dots, s^{(n)}, \theta) \\ &= \log \prod_{i=1}^n p_i(r^{(i)} | r^{(1)}, \dots, r^{(i-1)}, s^{(1)}, \dots, s^{(n)}, \theta) \\ &= \sum_{i=1}^n \log p_i(r^{(i)} | r^{(1)}, \dots, r^{(i-1)}, s^{(1)}, \dots, s^{(n)}, \theta)\end{aligned}$$

# The (log) likelihood

- For numerical reasons we work with  $\log p(\text{data}|\theta)$
- Using the rules of probability and logarithms:

$$\begin{aligned}\log p(\text{data}|\theta) &= \log p(r^{(1)}, \dots, r^{(n)} | s^{(1)}, \dots, s^{(n)}, \theta) \\ &= \log \prod_{i=1}^n p_i(r^{(i)} | r^{(1)}, \dots, r^{(i-1)}, s^{(1)}, \dots, s^{(n)}, \theta) \\ &= \sum_{i=1}^n \log p_i(r^{(i)} | r^{(1)}, \dots, r^{(i-1)}, s^{(1)}, \dots, s^{(n)}, \theta)\end{aligned}$$

- Simplest case:  $\log p(\text{data}|\theta) = \sum_{i=1}^n \log p_i(r^{(i)} | s^{(i)}, \theta)$

# The (log) likelihood

- For numerical reasons we work with  $\log p(\text{data}|\theta)$
- Using the rules of probability and logarithms:

$$\begin{aligned}\log p(\text{data}|\theta) &= \log p(r^{(1)}, \dots, r^{(n)} | s^{(1)}, \dots, s^{(n)}, \theta) \\ &= \log \prod_{i=1}^n p_i(r^{(i)} | r^{(1)}, \dots, r^{(i-1)}, s^{(1)}, \dots, s^{(n)}, \theta) \\ &= \sum_{i=1}^n \log p_i(r^{(i)} | r^{(1)}, \dots, r^{(i-1)}, s^{(1)}, \dots, s^{(n)}, \theta)\end{aligned}$$

- Simplest case:  $\log p(\text{data}|\theta) = \sum_{i=1}^n \log p_i(r^{(i)} | s^{(i)}, \theta)$
- Model building: Write function with
  - ▶ Input:  $\theta$  and data
  - ▶ Output:  $\log p(\text{data}|\theta)$

## Hacking time III

Let's play with a log-likelihood function

## 1 Introduction and recap

- Of models and likelihoods
- The psychometric function

## 2 Model fitting

- A statistical estimation problem
- Model fitting via point estimation

## 3 Intro to optimization

- The problem
- Optimization algorithms
- Optimization cheat sheets

Model fitting  $\sim$  statistical estimation problem

# Model fitting $\sim$ statistical estimation problem

## *Maximum likelihood estimation (MLE)*

- Find maximum of  $p(\text{data}|\theta)$

$$\hat{\theta}_{\text{ML}} = \arg \max_{\theta} p(\text{data}|\theta) = \arg \max_{\theta} \log p(\text{data}|\theta)$$

# Model fitting $\sim$ statistical estimation problem

## *Maximum likelihood estimation (MLE)*

- Find maximum of  $p(\text{data}|\theta)$

$$\hat{\theta}_{\text{ML}} = \arg \max_{\theta} p(\text{data}|\theta) = \arg \max_{\theta} \log p(\text{data}|\theta)$$

## Bayesian posterior estimation

$$p(\theta|\text{data}) = \frac{p(\text{data}|\theta)p(\theta)}{p(\text{data})} \propto p(\text{data}|\theta)p(\theta)$$

# Model fitting $\sim$ statistical estimation problem

## Maximum likelihood estimation (MLE)

- Find maximum of  $p(\text{data}|\theta)$

$$\hat{\theta}_{\text{ML}} = \arg \max_{\theta} p(\text{data}|\theta) = \arg \max_{\theta} \log p(\text{data}|\theta)$$

## Bayesian posterior estimation

$$p(\theta|\text{data}) = \frac{p(\text{data}|\theta)p(\theta)}{p(\text{data})} \propto p(\text{data}|\theta)p(\theta)$$

## Maximum-a-posteriori (MAP) estimation

$$\begin{aligned}\hat{\theta}_{\text{MAP}} &= \arg \max_{\theta} p(\theta|\text{data}) = \arg \max_{\theta} \log p(\theta|\text{data}) \\ &= \arg \max_{\theta} \{\log p(\text{data}|\theta) + \log p(\theta)\}\end{aligned}$$

# Two ways of model fitting

# Two ways of model fitting

## “Point estimation”

- Maximum likelihood estimation (MLE), Maximum-a-posteriori (MAP)
- Model fitting  $\sim$  *optimization problem*
- Returns a single ‘best’ parameter vector  $\hat{\theta}$  ( $\hat{\theta}_{ML}$  or  $\hat{\theta}_{MAP}$ )

# Two ways of model fitting

## "Point estimation"

- Maximum likelihood estimation (MLE), Maximum-a-posteriori (MAP)
- Model fitting  $\sim$  *optimization problem*
- Returns a single 'best' parameter vector  $\hat{\theta}$  ( $\hat{\theta}_{ML}$  or  $\hat{\theta}_{MAP}$ )

## Bayesian posterior estimation

- True Bayesian way
- Returns a *distribution* (the posterior distribution  $p(\theta|\text{data})$ )
- For  $n \rightarrow \infty$  converges to MLE (if  $p(\hat{\theta}_{ML}) \neq 0$ )
- In practice, it returns an *approximation* of the posterior

# Two ways of model fitting

## "Point estimation"

- Maximum likelihood estimation (MLE), Maximum-a-posteriori (MAP)
- Model fitting  $\sim$  *optimization problem*
- Returns a single 'best' parameter vector  $\hat{\theta}$  ( $\hat{\theta}_{ML}$  or  $\hat{\theta}_{MAP}$ )

## Bayesian posterior estimation

- True Bayesian way
- Returns a *distribution* (the posterior distribution  $p(\theta|\text{data})$ )
- For  $n \rightarrow \infty$  converges to MLE (if  $p(\hat{\theta}_{ML}) \neq 0$ )
- In practice, it returns an *approximation* of the posterior
  - ① In parametric form  $q(\theta)$  (e.g., *variational inference*)
  - ② As a bunch of discrete samples (e.g., *Markov-Chain Monte Carlo*)

# Two ways of model fitting

## "Point estimation"

- Maximum likelihood estimation (MLE), Maximum-a-posteriori (MAP)
- Model fitting  $\sim$  *optimization problem*
- Returns a single 'best' parameter vector  $\hat{\theta}$  ( $\hat{\theta}_{ML}$  or  $\hat{\theta}_{MAP}$ )

## Bayesian posterior estimation

- True Bayesian way
- Returns a *distribution* (the posterior distribution  $p(\theta|\text{data})$ )
- For  $n \rightarrow \infty$  converges to MLE (if  $p(\hat{\theta}_{ML}) \neq 0$ )
- In practice, it returns an *approximation* of the posterior
  - ① In parametric form  $q(\theta)$  (e.g., *variational inference*)
  - ② As a bunch of discrete samples (e.g., *Markov-Chain Monte Carlo*)
- Next lecture!

## Model fitting via point estimation

- Find single  $\theta$  that best describes the data

## Model fitting via point estimation

- Find single  $\theta$  that best describes the data
- (For this section we switch notation from  $\theta$  to  $x$ )

## Model fitting via point estimation

- Find single  $\theta$  that best describes the data
- (For this section we switch notation from  $\theta$  to  $x$ )
- Given  $\tilde{f}(x) \equiv \begin{cases} \log p(\text{data}|x) & \text{maximum likelihood} \\ \log p(\text{data}|x) + \log p(x) & \text{maximum-a-posteriori} \end{cases}$

## Model fitting via point estimation

- Find single  $\theta$  that best describes the data
- (For this section we switch notation from  $\theta$  to  $x$ )
- Given  $\tilde{f}(x) \equiv \begin{cases} \log p(\text{data}|x) & \text{maximum likelihood} \\ \log p(\text{data}|x) + \log p(x) & \text{maximum-a-posteriori} \end{cases}$
- By convention, we *minimize*  $f(x) \equiv -\tilde{f}(x)$

# Model fitting via point estimation

- Find single  $\theta$  that best describes the data
- (For this section we switch notation from  $\theta$  to  $x$ )
- Given  $\tilde{f}(x) \equiv \begin{cases} \log p(\text{data}|x) & \text{maximum likelihood} \\ \log p(\text{data}|x) + \log p(x) & \text{maximum-a-posteriori} \end{cases}$
- By convention, we minimize  $f(x) \equiv -\tilde{f}(x)$
- $\Rightarrow$  Find  $x_{opt} \approx \arg \min_x f(x)$  as fast as possible

# Model fitting via point estimation

- Find single  $\theta$  that best describes the data
- (For this section we switch notation from  $\theta$  to  $x$ )
- Given  $\tilde{f}(x) \equiv \begin{cases} \log p(\text{data}|x) & \text{maximum likelihood} \\ \log p(\text{data}|x) + \log p(x) & \text{maximum-a-posteriori} \end{cases}$
- By convention, we minimize  $f(x) \equiv -\tilde{f}(x)$
- $\Rightarrow$  Find  $x_{opt} \approx \arg \min_x f(x)$  as fast as possible
- General case:  $f(x)$  is a *black box*
  - ▶ Sometimes we can compute the gradient

# Model fitting via point estimation

- Find single  $\theta$  that best describes the data
- (For this section we switch notation from  $\theta$  to  $x$ )
- Given  $\tilde{f}(x) \equiv \begin{cases} \log p(\text{data}|x) & \text{maximum likelihood} \\ \log p(\text{data}|x) + \log p(x) & \text{maximum-a-posteriori} \end{cases}$
- By convention, we minimize  $f(x) \equiv -\tilde{f}(x)$
- $\Rightarrow$  Find  $x_{opt} \approx \arg \min_x f(x)$  as fast as possible
- General case:  $f(x)$  is a *black box*
  - ▶ Sometimes we can compute the gradient

General solution

Feed  $f(x)$  to an optimization algorithm

## 1 Introduction and recap

- Of models and likelihoods
- The psychometric function

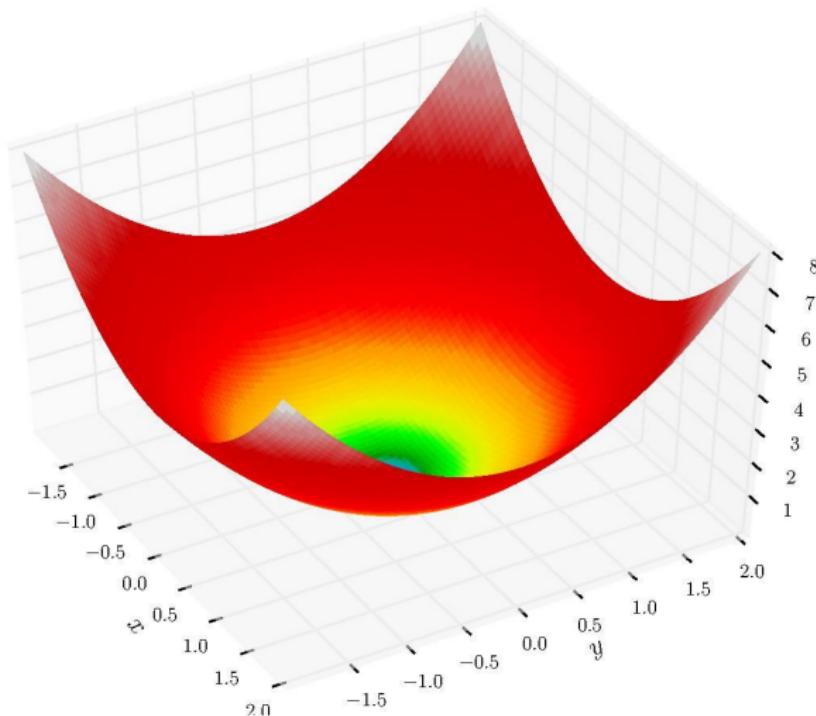
## 2 Model fitting

- A statistical estimation problem
- Model fitting via point estimation

## 3 Intro to optimization

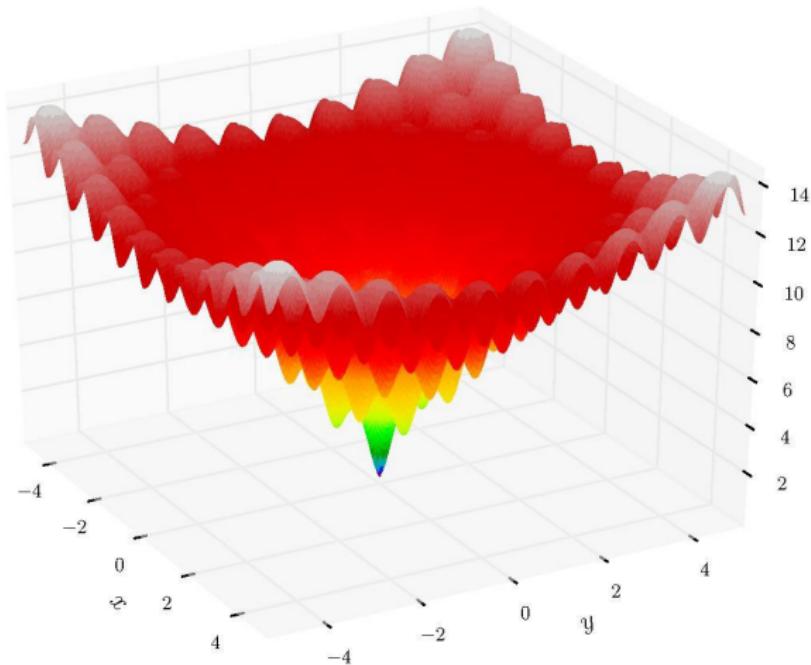
- The problem
- Optimization algorithms
- Optimization cheat sheets

# How hard can it be?



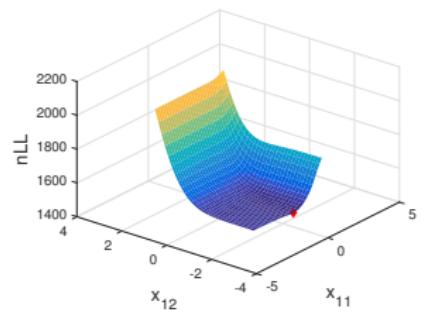
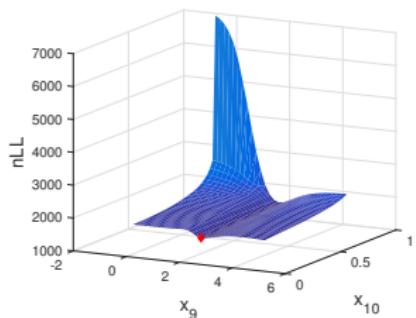
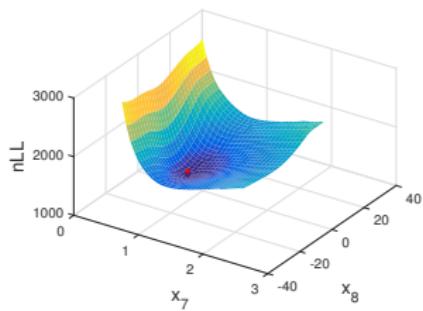
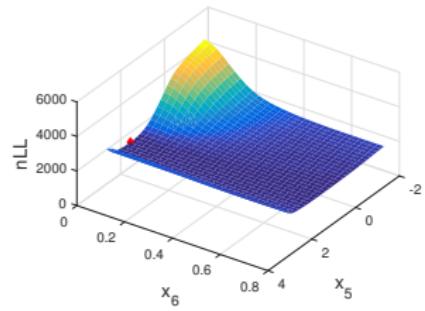
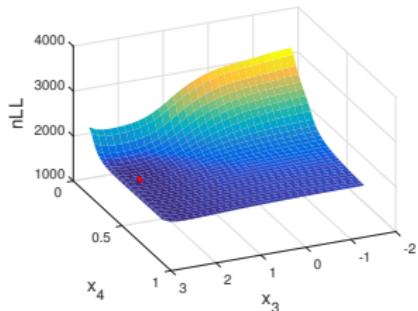
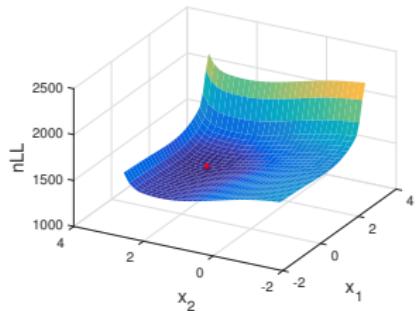
Source: Wikimedia Commons

# How hard can it be?

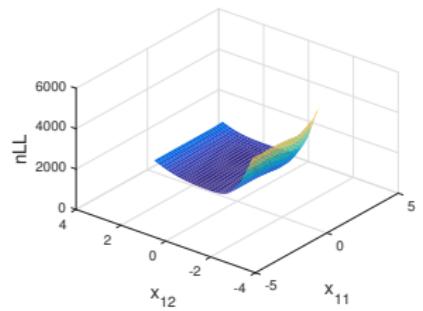
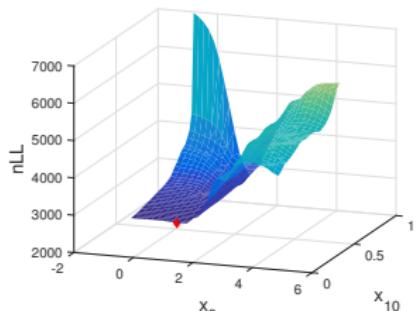
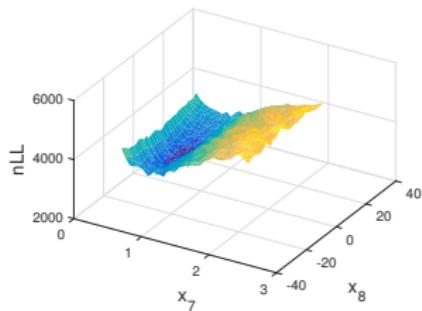
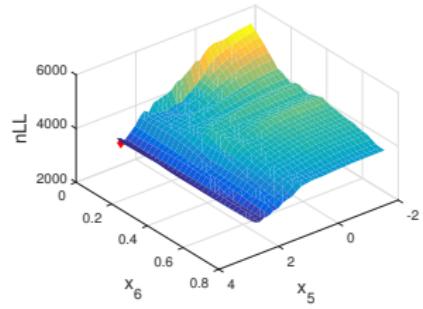
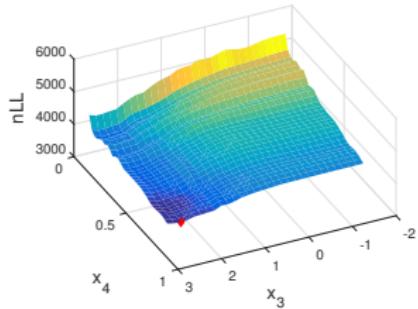
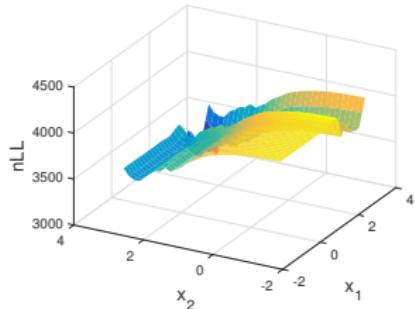


Source: Wikimedia Commons

# How hard can it be?



# How hard can it be?



# How hard can it be?

neval	$x_1$	$x_2$	$f(x)$
1	-0.500	2.500	508.500
2	-0.525	2.500	497.110
3	-0.500	2.625	566.313
4	-0.525	2.375	443.063
5	-0.537	2.250	386.953
6	-0.563	2.250	376.320
7	-0.594	2.125	316.702
8	-0.606	1.875	229.824
9	-0.647	1.563	133.598
10	-0.703	1.438	91.847
11	-0.786	1.031	20.292
12	-0.839	0.469	8.918
13	-0.962	-0.359	168.785
14	-0.978	-0.063	107.796
15	-0.895	0.344	24.553
16	-0.730	1.156	41.905
17	-0.854	0.547	6.760
18	-0.907	-0.016	73.917
19	-0.816	0.770	4.366
20	-0.831	0.848	5.818
21	-0.793	1.070	22.655
22	-0.839	0.678	3.448
23	-0.824	0.600	3.955
24	-0.846	0.508	7.766
25	-0.824	0.704	3.391
26	-0.839	0.782	4.004
27	-0.828	0.645	3.497
28	-0.835	0.737	3.523
29	?	?	?

# Optimization can be hard

- ➊ Optimizer does not see the landscape!

# Optimization can be hard

- ① Optimizer does not see the landscape!
- ② Multiple local minima or saddle points ('non-convex')

# Optimization can be hard

- ① Optimizer does not see the landscape!
- ② Multiple local minima or saddle points ('non-convex')
- ③ Expensive function evaluation

# Optimization can be hard

- ① Optimizer does not see the landscape!
- ② Multiple local minima or saddle points ('non-convex')
- ③ Expensive function evaluation
- ④ Noisy function evaluation
- ⑤ Rough landscape (numerical approximations, etc.)

# Optimization algorithms

## Gradient-based methods

- Stochastic gradient descent (e.g., ADAM)
- Quasi-Newton methods (e.g., BFGS aka `fminunc/fmincon`)

## Gradient-free methods

- Nelder-Mead (`fminsearch`)
- Pattern/direct search (`patternsearch`)
- Simulated annealing
- Genetic algorithms
- CMA-ES
- Bayesian optimization
- Bayesian Adaptive Direct Search (BADS; Acerbi & Ma, *NeurIPS* 2017)

# Optimization algorithms

## Gradient-based methods

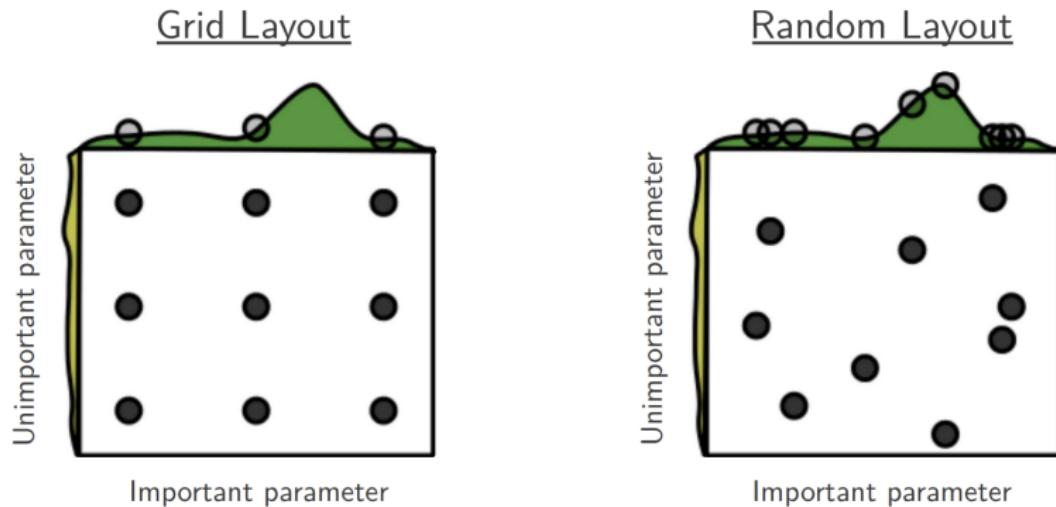
- Stochastic gradient descent (e.g., ADAM)
- Quasi-Newton methods (e.g., BFGS aka `fminunc/fmincon`)

## Gradient-free methods

- Nelder-Mead (`fminsearch`)
- Pattern/direct search (`patternsearch`)
- Simulated annealing
- Genetic algorithms
- CMA-ES
- Bayesian optimization
- Bayesian Adaptive Direct Search (BADS; Acerbi & Ma, *NeurIPS* 2017)

Demos: <https://github.com/lacerbi/optimviz>

# Grid and random search

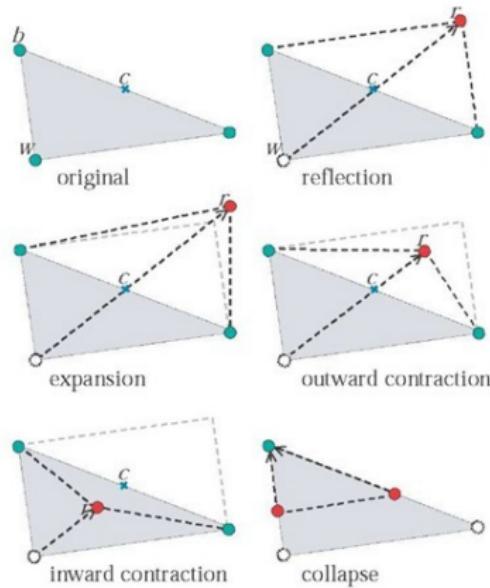


Source: Bergstra & Bengio (2012)

- Curse of dimensionality
- $N$  points per (important) dimension  $\Rightarrow N^D$  points

# Nelder-Mead (fminsearch)

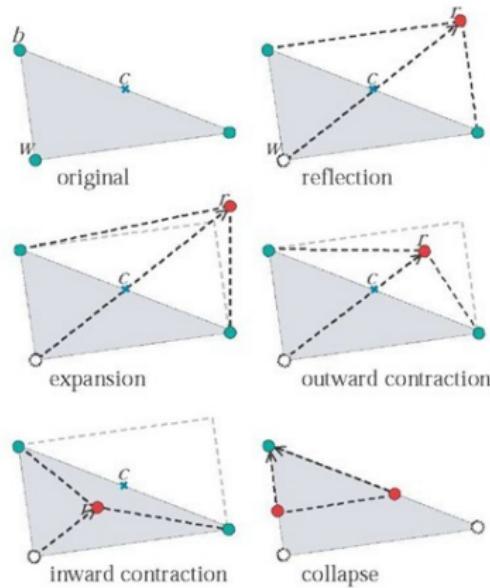
J. A. Nelder & R. Mead, A simplex method for function minimization (1965)



Source: Encyclopedia of Artificial Intelligence (2009)

# Nelder-Mead (fminsearch)

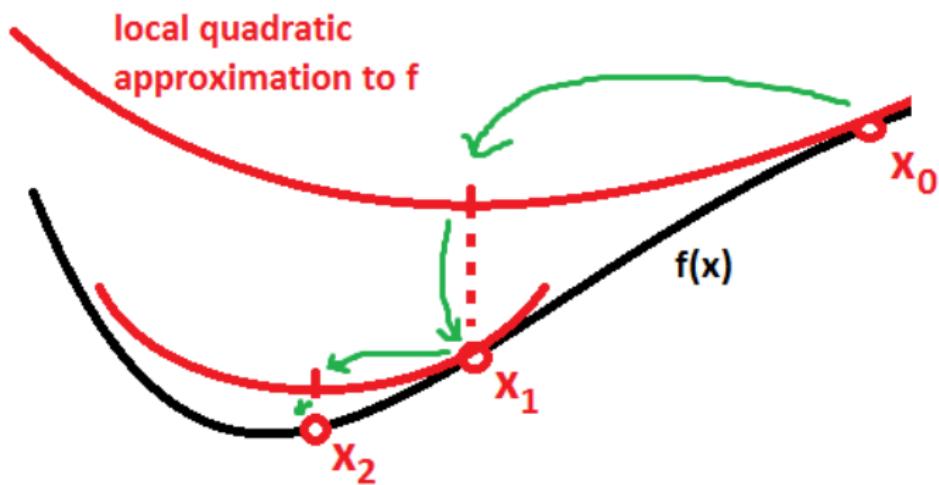
J. A. Nelder & R. Mead, A simplex method for function minimization (1965)



Source: Encyclopedia of Artificial Intelligence (2009)

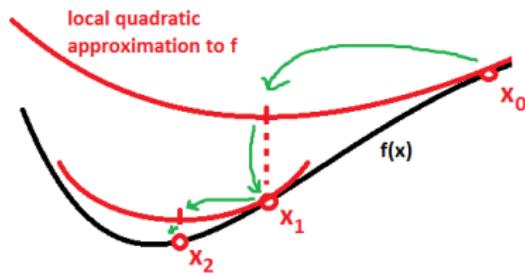
Bounded optimization: **fminsearchbnd** (John d'Errico)

## Newton method



Source: StackExchange

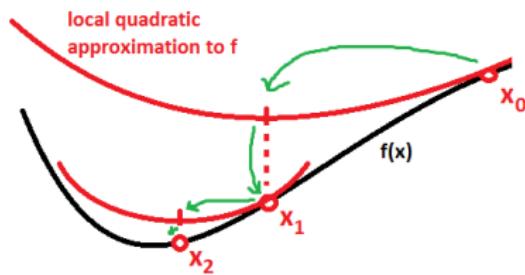
# Newton method



Source: StackExchange

Needs the inverse of the curvature (inverse Hessian)  
Very expensive in high dimension

## Quasi-Newton methods (fminunc, fmincon)

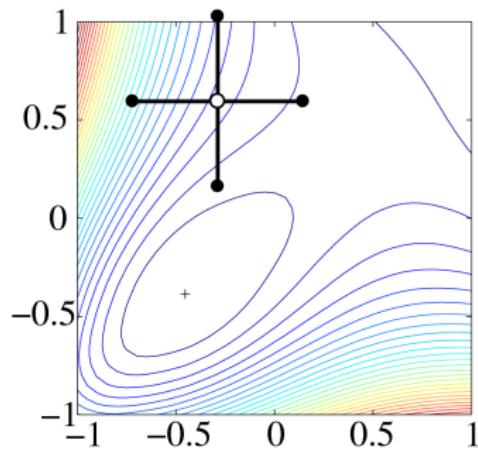


Source: StackExchange

Approximate Hessian (DFP) or inverse Hessian (BFGS) via gradient  
Very fast and efficient on smooth problems

## Direct search (patternsearch)

R. Hooke and T.A. Jeeves, "Direct search" solution of numerical and statistical problems (1961)



Source: Wikimedia Commons

# Genetic Algorithms (ga)

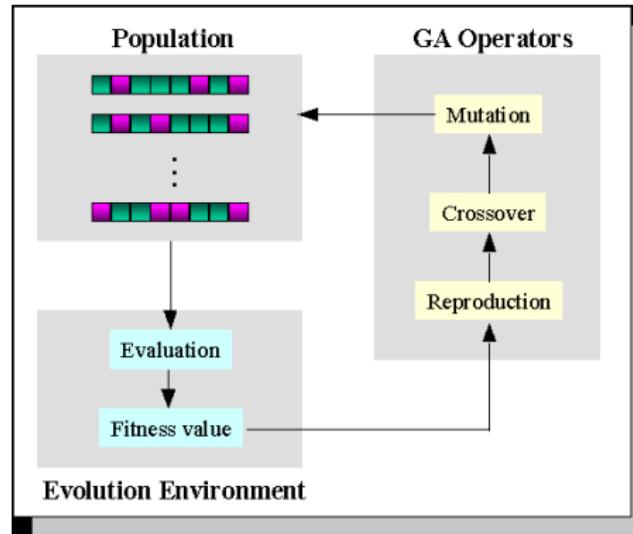
J.H. Holland, Adaptation in Natural and Artificial Systems (1975)

- Evolutionary algorithm
- Population based

# Genetic Algorithms (ga)

J.H. Holland, Adaptation in Natural and Artificial Systems (1975)

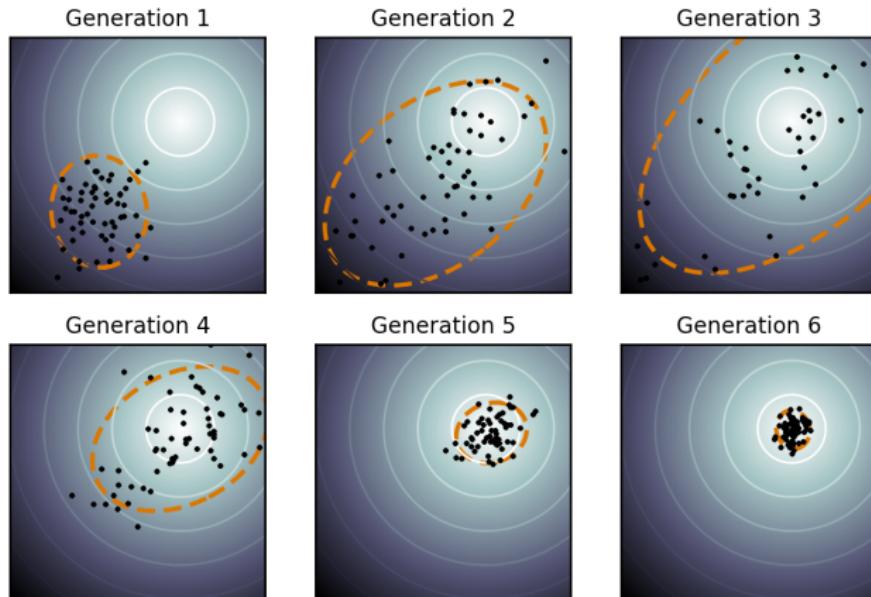
- Evolutionary algorithm
- Population based



Source: An Educational GA Learning Tool (IEEE)

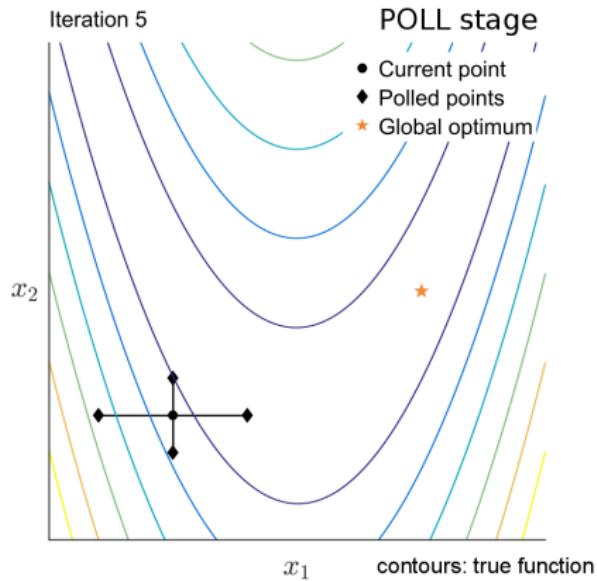
# Cov. Matrix Adaptation - Evolution Strategies (cmaes)

[\*] N. Hansen, S. D. Müller, P. Koumoutsakos, Reducing the time complexity of the derandomized evolution strategy with covariance matrix adaptation (CMA-ES), (2003)



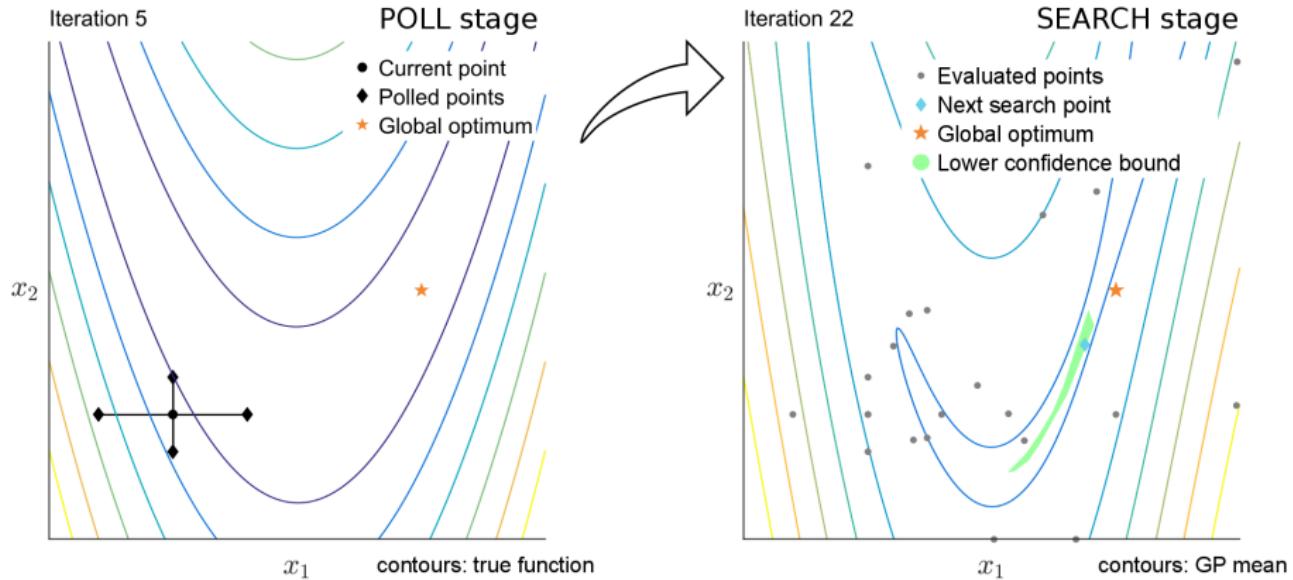
Source: Wikipedia

# Bayesian Adaptive Direct Search (bads)



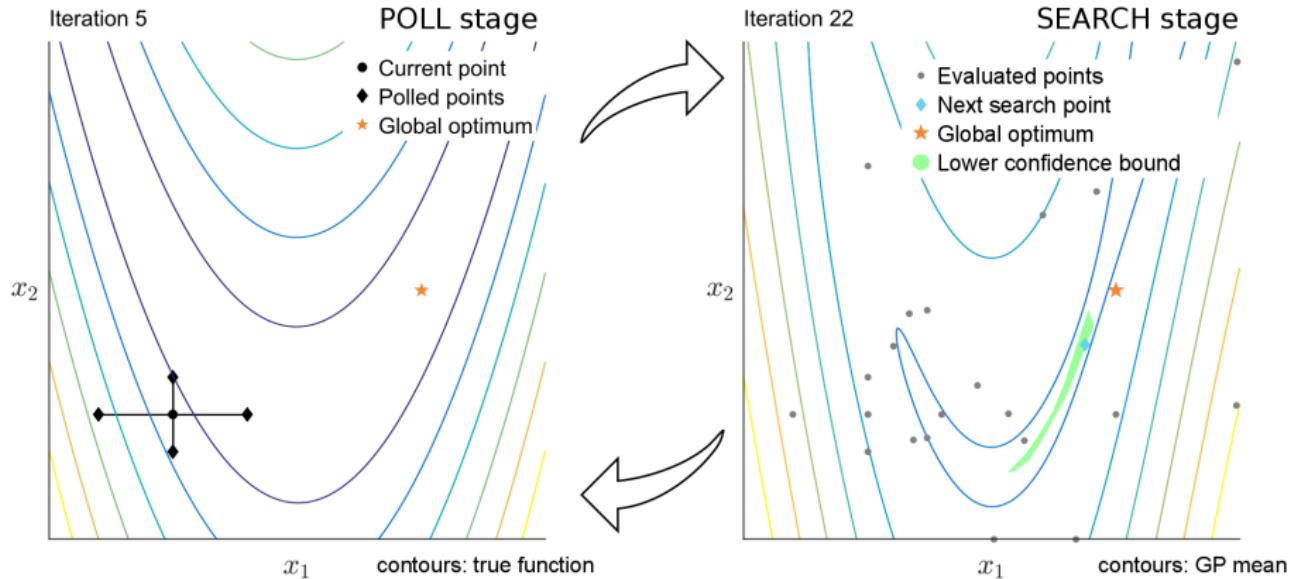
Acerbi & Ma, *NeurIPS* (2017)

# Bayesian Adaptive Direct Search (bads)



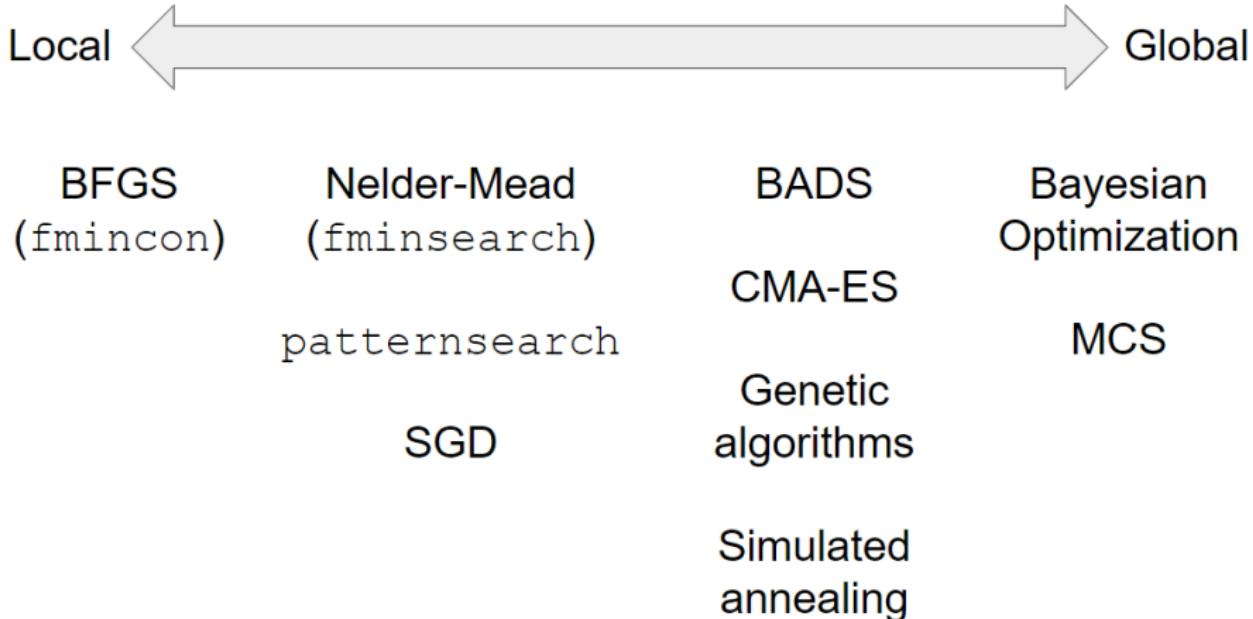
Acerbi & Ma, *NeurIPS* (2017)

# Bayesian Adaptive Direct Search (bads)



Acerbi & Ma, *NeurIPS* (2017)

# Local vs. global optimization



# Preparing for optimization

- *Domain* of parameter vector  $\boldsymbol{\theta} = (\theta_1, \theta_2, \dots, \theta_D) \in \Theta$

# Preparing for optimization

- Domain of parameter vector  $\theta = (\theta_1, \theta_2, \dots, \theta_D) \in \Theta$

In practice, for each  $\theta_d$ , define

- ▶ The *hard bounds* of the parameter.
  - ★ Mathematical constraints (e.g.,  $\sigma > 0$ ;  $0 \leq p \leq 1$ )
  - ★ Effective physical limitations

# Preparing for optimization

- Domain of parameter vector  $\theta = (\theta_1, \theta_2, \dots, \theta_D) \in \Theta$

In practice, for each  $\theta_d$ , define

- ▶ The *hard bounds* of the parameter.
  - ★ Mathematical constraints (e.g.,  $\sigma > 0$ ;  $0 \leq p \leq 1$ )
  - ★ Effective physical limitations
- ▶ The *plausible bounds* of the parameter
  - ★ Should span parameter values for most datasets (e.g., 95% prior interval)
  - ★ Built from pilot studies, literature, guesswork
  - ★ If in doubt, start larger

# Preparing for optimization

- Domain of parameter vector  $\theta = (\theta_1, \theta_2, \dots, \theta_D) \in \Theta$

In practice, for each  $\theta_d$ , define

- ▶ The *hard bounds* of the parameter.
  - ★ Mathematical constraints (e.g.,  $\sigma > 0$ ;  $0 \leq p \leq 1$ )
  - ★ Effective physical limitations
- ▶ The *plausible bounds* of the parameter
  - ★ Should span parameter values for most datasets (e.g., 95% prior interval)
  - ★ Built from pilot studies, literature, guesswork
  - ★ If in doubt, start larger
- Consider reparameterizations to achieve
  - ▶ Uniformity of effects across parameter range
  - ▶ Independence between parameters

## Hacking time IV

Let's optimize the log-likelihood for the psychometric model

# Optimization cheat sheet, page 1

## Rule zero

Understand your problem  $\implies$  often a *gray box*

# Optimization cheat sheet, page 1

## Rule zero

Understand your problem  $\implies$  often a *gray box*

### Input variables:

- Dimensionality: low ( $D \lesssim 10$ ) or high ( $D \gg 20$ )

# Optimization cheat sheet, page 1

## Rule zero

Understand your problem  $\implies$  often a *gray box*

### Input variables:

- Dimensionality: low ( $D \lesssim 10$ ) or high ( $D \gg 20$ )
- Bounds: Think of *hard* and *plausible* bounds

# Optimization cheat sheet, page 1

## Rule zero

Understand your problem  $\implies$  often a *gray box*

### Input variables:

- Dimensionality: low ( $D \lesssim 10$ ) or high ( $D \gg 20$ )
- Bounds: Think of *hard* and *plausible* bounds
- Parameterization: Not all parameterizations are created equal

# Optimization cheat sheet, page 1

## Rule zero

Understand your problem  $\implies$  often a gray box

### Input variables:

- Dimensionality: low ( $D \lesssim 10$ ) or high ( $D \gg 20$ )
- Bounds: Think of *hard* and *plausible* bounds
- Parameterization: Not all parameterizations are created equal

### Target function:

- Convexity: convex or non-convex

# Optimization cheat sheet, page 1

## Rule zero

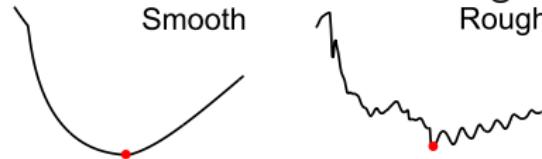
Understand your problem  $\Rightarrow$  often a gray box

### Input variables:

- Dimensionality: low ( $D \lesssim 10$ ) or high ( $D \gg 20$ )
- Bounds: Think of *hard* and *plausible* bounds
- Parameterization: Not all parameterizations are created equal

### Target function:

- Convexity: convex or non-convex
- Smoothness: smooth or rough



# Optimization cheat sheet, page 1

## Rule zero

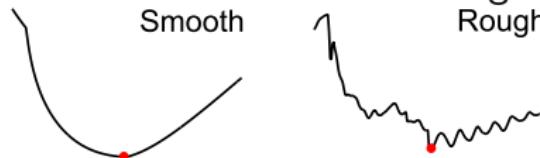
Understand your problem  $\Rightarrow$  often a gray box

### Input variables:

- Dimensionality: low ( $D \lesssim 10$ ) or high ( $D \gg 20$ )
- Bounds: Think of *hard* and *plausible* bounds
- Parameterization: Not all parameterizations are created equal

### Target function:

- Convexity: convex or non-convex
- Smoothness: smooth or rough



- Deterministic or stochastic

# Optimization cheat sheet, page 1

## Rule zero

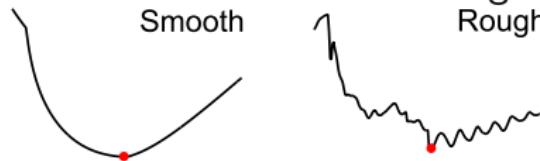
Understand your problem  $\Rightarrow$  often a gray box

### Input variables:

- Dimensionality: low ( $D \lesssim 10$ ) or high ( $D \gg 20$ )
- Bounds: Think of *hard* and *plausible* bounds
- Parameterization: Not all parameterizations are created equal

### Target function:

- Convexity: convex or non-convex
- Smoothness: smooth or rough



- Deterministic or stochastic
  - ▶ If stochastic  $\Rightarrow$  minimize  $\mathbb{E}[f(x)]$

# Optimization cheat sheet, page 1

## Rule zero

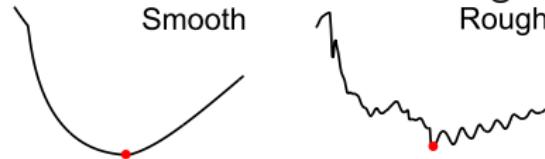
Understand your problem  $\Rightarrow$  often a gray box

### Input variables:

- Dimensionality: low ( $D \lesssim 10$ ) or high ( $D \gg 20$ )
- Bounds: Think of *hard* and *plausible* bounds
- Parameterization: Not all parameterizations are created equal

### Target function:

- Convexity: convex or non-convex
- Smoothness: smooth or rough



- Deterministic or stochastic
  - ▶ If stochastic  $\Rightarrow$  minimize  $\mathbb{E}[f(x)]$
- Computational cost:  
cheap ( $\ll 0.01$  s), moderate (0.01-1 s), or expensive ( $\gg 1$  s)

# Optimization cheat sheet, page 2

## Fundamental theorem

'No Free Lunch' theorem  $\implies$  no single best optimizer for all problems

# Optimization cheat sheet, page 2

## Fundamental theorem

'No Free Lunch' theorem  $\implies$  no single best optimizer for all problems  
(But not all methods are created equal!)

# Optimization cheat sheet, page 2

## Fundamental theorem

'No Free Lunch' theorem  $\implies$  no single best optimizer for all problems  
(But not all methods are created equal!)

- Is your problem smooth?

# Optimization cheat sheet, page 2

## Fundamental theorem

'No Free Lunch' theorem  $\implies$  no single best optimizer for all problems  
(But not all methods are created equal!)

- Is your problem smooth?
  - ▶ If you have the gradient  $\implies$  BFGS

# Optimization cheat sheet, page 2

## Fundamental theorem

'No Free Lunch' theorem  $\implies$  no single best optimizer for all problems  
(But not all methods are created equal!)

- Is your problem smooth?
  - ▶ If you have the gradient  $\implies$  BFGS
  - ▶ If low- $D$  and cheap  $\implies$  BFGS with finite differences

# Optimization cheat sheet, page 2

## Fundamental theorem

'No Free Lunch' theorem  $\implies$  no single best optimizer for all problems  
(But not all methods are created equal!)

- Is your problem smooth?
  - ▶ If you have the gradient  $\implies$  BFGS
  - ▶ If low- $D$  and cheap  $\implies$  BFGS with finite differences
  - ▶ If low- $D$  and (moderately) costly  $\implies$  BADS

# Optimization cheat sheet, page 2

## Fundamental theorem

'No Free Lunch' theorem  $\implies$  no single best optimizer for all problems  
(But not all methods are created equal!)

- Is your problem smooth?
  - ▶ If you have the gradient  $\implies$  BFGS
  - ▶ If low- $D$  and cheap  $\implies$  BFGS with finite differences
  - ▶ If low- $D$  and (moderately) costly  $\implies$  BADS
- Is your problem rough or noisy?

# Optimization cheat sheet, page 2

## Fundamental theorem

'No Free Lunch' theorem  $\implies$  no single best optimizer for all problems  
(But not all methods are created equal!)

- Is your problem smooth?
  - ▶ If you have the gradient  $\implies$  BFGS
  - ▶ If low- $D$  and cheap  $\implies$  BFGS with finite differences
  - ▶ If low- $D$  and (moderately) costly  $\implies$  BADS
- Is your problem rough or noisy?
  - ▶ First, try and make it smooth and deterministic!

# Optimization cheat sheet, page 2

## Fundamental theorem

'No Free Lunch' theorem  $\implies$  no single best optimizer for all problems  
(But not all methods are created equal!)

- Is your problem smooth?
  - ▶ If you have the gradient  $\implies$  BFGS
  - ▶ If low- $D$  and cheap  $\implies$  BFGS with finite differences
  - ▶ If low- $D$  and (moderately) costly  $\implies$  BADS
- Is your problem rough or noisy?
  - ▶ First, try and make it smooth and deterministic!
  - ▶ If gradient is available and high- $D$   $\implies$  SGD (e.g., ADAM)

# Optimization cheat sheet, page 2

## Fundamental theorem

'No Free Lunch' theorem  $\implies$  no single best optimizer for all problems  
(But not all methods are created equal!)

- Is your problem smooth?
  - ▶ If you have the gradient  $\implies$  BFGS
  - ▶ If low- $D$  and cheap  $\implies$  BFGS with finite differences
  - ▶ If low- $D$  and (moderately) costly  $\implies$  BADS
- Is your problem rough or noisy?
  - ▶ First, try and make it smooth and deterministic!
  - ▶ If gradient is available and high- $D$   $\implies$  SGD (e.g., ADAM)
  - ▶ If high- $D$  and cheap  $\implies$  CMA-ES

# Optimization cheat sheet, page 2

## Fundamental theorem

'No Free Lunch' theorem  $\implies$  no single best optimizer for all problems  
(But not all methods are created equal!)

- Is your problem smooth?
  - ▶ If you have the gradient  $\implies$  BFGS
  - ▶ If low- $D$  and cheap  $\implies$  BFGS with finite differences
  - ▶ If low- $D$  and (moderately) costly  $\implies$  BADS
- Is your problem rough or noisy?
  - ▶ First, try and make it smooth and deterministic!
  - ▶ If gradient is available and high- $D$   $\implies$  SGD (e.g., ADAM)
  - ▶ If high- $D$  and cheap  $\implies$  CMA-ES
  - ▶ If low- $D$  and (moderately) costly  $\implies$  BADS

# Optimization cheat sheet, page 2

## Fundamental theorem

'No Free Lunch' theorem  $\implies$  no single best optimizer for all problems  
(But not all methods are created equal!)

- Is your problem smooth?
  - ▶ If you have the gradient  $\implies$  BFGS
  - ▶ If low- $D$  and cheap  $\implies$  BFGS with finite differences
  - ▶ If low- $D$  and (moderately) costly  $\implies$  BADS
- Is your problem rough or noisy?
  - ▶ First, try and make it smooth and deterministic!
  - ▶ If gradient is available and high- $D$   $\implies$  SGD (e.g., ADAM)
  - ▶ If high- $D$  and cheap  $\implies$  CMA-ES
  - ▶ If low- $D$  and (moderately) costly  $\implies$  BADS
- Is your problem high- $D$ , costly, and you do not have the gradient?

# Optimization cheat sheet, page 2

## Fundamental theorem

'No Free Lunch' theorem  $\implies$  no single best optimizer for all problems  
(But not all methods are created equal!)

- Is your problem smooth?
  - ▶ If you have the gradient  $\implies$  BFGS
  - ▶ If low- $D$  and cheap  $\implies$  BFGS with finite differences
  - ▶ If low- $D$  and (moderately) costly  $\implies$  BADS
- Is your problem rough or noisy?
  - ▶ First, try and make it smooth and deterministic!
  - ▶ If gradient is available and high- $D$   $\implies$  SGD (e.g., ADAM)
  - ▶ If high- $D$  and cheap  $\implies$  CMA-ES
  - ▶ If low- $D$  and (moderately) costly  $\implies$  BADS
- Is your problem high- $D$ , costly, and you do not have the gradient?
  - ▶ Tough luck

# Optimization cheat sheet, page 3

## The golden rule

No optimizer can guarantee to find the global optimum

# Optimization cheat sheet, page 3

## The golden rule

No optimizer can guarantee to find the global optimum

⇒ **Always** perform multiple distinct optimization runs ('restarts')

# Optimization cheat sheet, page 3

## The golden rule

No optimizer can guarantee to find the global optimum

⇒ **Always** perform multiple distinct optimization runs ('restarts')

- How to choose starting points?

# Optimization cheat sheet, page 3

## The golden rule

No optimizer can guarantee to find the global optimum

⇒ **Always** perform multiple distinct optimization runs ('restarts')

- How to choose starting points?
  - ▶ Draw from prior distribution

# Optimization cheat sheet, page 3

## The golden rule

No optimizer can guarantee to find the global optimum

⇒ **Always** perform multiple distinct optimization runs ('restarts')

- How to choose starting points?
  - ▶ Draw from prior distribution
  - ▶ Draw from a 'plausible' box

# Optimization cheat sheet, page 3

## The golden rule

No optimizer can guarantee to find the global optimum

⇒ **Always** perform multiple distinct optimization runs ('restarts')

- How to choose starting points?
  - ▶ Draw from prior distribution
  - ▶ Draw from a 'plausible' box
  - ▶ Sieve method

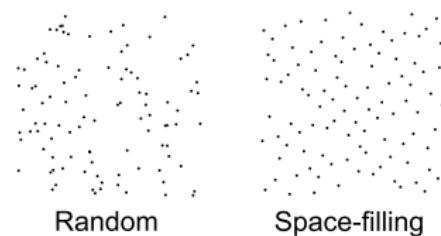
# Optimization cheat sheet, page 3

## The golden rule

No optimizer can guarantee to find the global optimum

⇒ **Always** perform multiple distinct optimization runs ('restarts')

- How to choose starting points?
  - ▶ Draw from prior distribution
  - ▶ Draw from a 'plausible' box
  - ▶ Sieve method
  - ▶ Use space-filling designs  
(quasi-random sequences)



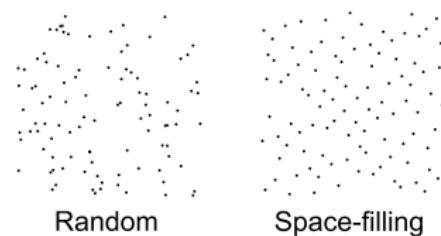
# Optimization cheat sheet, page 3

## The golden rule

No optimizer can guarantee to find the global optimum

⇒ **Always** perform multiple distinct optimization runs ('restarts')

- How to choose starting points?
  - ▶ Draw from prior distribution
  - ▶ Draw from a 'plausible' box
  - ▶ Sieve method
  - ▶ Use space-filling designs  
(quasi-random sequences)
  
- How many restarts?



# Optimization cheat sheet, page 3

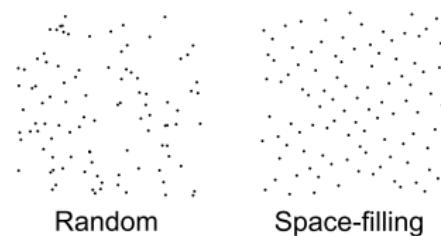
## The golden rule

No optimizer can guarantee to find the global optimum

⇒ **Always** perform multiple distinct optimization runs ('restarts')

- How to choose starting points?

- ▶ Draw from prior distribution
- ▶ Draw from a 'plausible' box
- ▶ Sieve method
- ▶ Use space-filling designs  
(quasi-random sequences)



- How many restarts?

- ▶ As many as you need

# Optimization cheat sheet, page 3

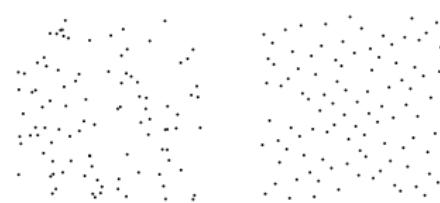
## The golden rule

No optimizer can guarantee to find the global optimum

⇒ **Always** perform multiple distinct optimization runs ('restarts')

- How to choose starting points?

- ▶ Draw from prior distribution
- ▶ Draw from a 'plausible' box
- ▶ Sieve method
- ▶ Use space-filling designs  
(quasi-random sequences)



- How many restarts?

- ▶ As many as you need
- ▶ Informally, check that 'most' points converge to the same solution

# Optimization cheat sheet, page 3

## The golden rule

No optimizer can guarantee to find the global optimum

⇒ **Always** perform multiple distinct optimization runs ('restarts')

- How to choose starting points?

- ▶ Draw from prior distribution
- ▶ Draw from a 'plausible' box
- ▶ Sieve method
- ▶ Use space-filling designs  
(quasi-random sequences)



- How many restarts?

- ▶ As many as you need
- ▶ Informally, check that 'most' points converge to the same solution
- ▶ *Bootstrap* approach (Acerbi, Dokka et al., *PLoS Comp Biol* 2018)

# Learning objectives (did I lie?)

By the end of this lecture/tutorial, we will:

- Explain how model fitting is an **optimization problem**
- Describe several major **optimization algorithms**
- Set up and run optimization on **real dataset and models**
- Summarize practical **tips & tricks** for optimization

# Final slide

- Contacts: `luigi.acerbi@helsinki.fi`, @AcerbiLuigi

## Code:

- BADS (MATLAB): `github.com/lacerbi/bads`
- PyBADS (Python): In the works!
- Optimization demos: `github.com/lacerbi/optimviz`

# Final slide

- Contacts: `luigi.acerbi@helsinki.fi`, @AcerbiLuigi

## Code:

- BADS (MATLAB): `github.com/lacerbi/bads`
- PyBADS (Python): In the works!
- Optimization demos: `github.com/lacerbi/optimviz`

Thanks!

(Time for questions?)